



Pig Latin: A Not-So-Foreign Language for Data Processing



Motivation

“You”re a procedural programmer

“You have huge data

“You want to analyze it

Motivation

“ As a procedural programmer...

“ May find writing queries in SQL unnatural and too restrictive

“ More comfortable with writing code; a series of statements as opposed to a long query. (Ex: MapReduce is so successful).

Motivation

“ Data analysis goals

“ Quick

- “ Exploit parallel processing power of a distributed system

“ Easy

- “ Be able to write a program or query without a huge learning curve
- “ Have some common analysis tasks predefined

“ Flexible

- “ Transform a data set(s) into a workable structure without much overhead
- “ Perform customized processing

“ Transparent

- “ Have a say in how the data processing is executed on the system

Motivation

“ Relational Distributed Databases

- “ Parallel database products expensive
- “ Rigid schemas
- “ Processing requires declarative SQL query construction

“ Map-Reduce

- “ Relies on custom code for even common operations
- “ Need to do workarounds for tasks that have different data flows other than the expected Map→Combine→Reduce

Motivation

“ Relational Distributed Databases

“ Sweet Spot: Take the best of both SQL and Map-Reduce; combine high-level declarative querying with low-level procedural programming...Pig Latin!

“ Map-Reduce

Pig Latin Example

Table urls: (url,category, pagerank)

Find for each sufficiently large category, the average pagerank of high-pagerank urls in that category

SQL:

```
SELECT category,AVG(pagerank)
FROM urls WHERE pagerank > 0.2
GROUP BY category HAVING COUNT(*) > 10^6
```

Pig Latin:

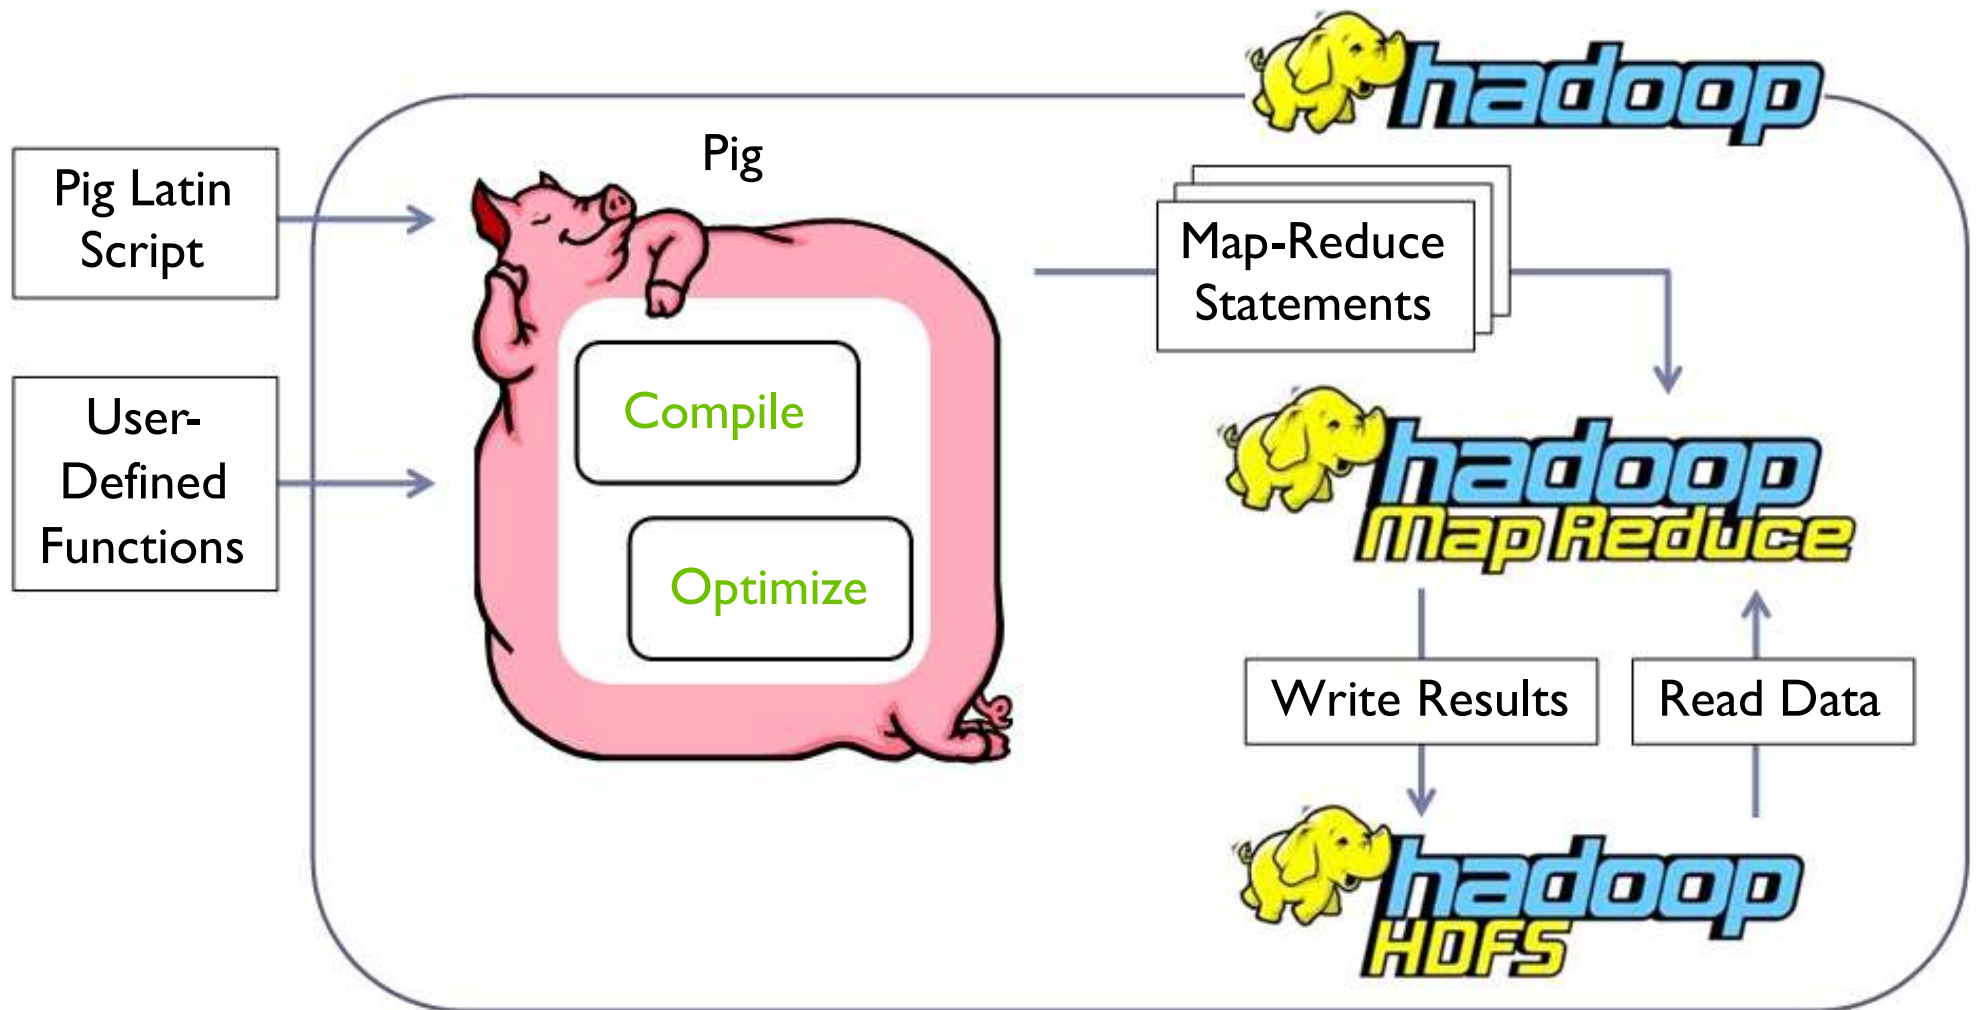
```
good_urls = FILTER urls BY pagerank > 0.2;
groups = GROUP good_urls BY category;
big_groups = FILTER groups BY COUNT(good_urls)>10^6;
output = FOREACH big_groups GENERATE category,
AVG(good_urls.pagerank);
```



Outline

- “ System Overview
- “ Pig Latin (The Language)
 - “ Data Structures
 - “ Commands
- “ Pig (The Compiler)
 - “ Logical & Physical Plans
 - “ Optimization
 - “ Efficiency
- “ Pig Pen (The Debugger)
- “ Conclusion

Big Picture



Data Model

- “ Atom - simple atomic value (ie: number or string)
- “ Tuple
- “ Bag
- “ Map

$$\left(\text{'alice'}, \left\{ \begin{array}{l} (\text{'lakers'}, 1) \\ (\text{'iPod'}, 2) \end{array} \right\}, [\text{'age'} \rightarrow 20] \right)$$

Data Model

- “ Atom
- “ Tuple - sequence of fields; each field any type
- “ Bag
- “ Map

$$\left(\text{'alice'}, \left\{ \begin{array}{l} \text{'lakers', 1} \\ \text{'iPod', 2} \end{array} \right\}, [\text{'age'} \rightarrow 20] \right)$$

Data Model

“ Atom

“ Tuple

“ Bag - collection of tuples

“ Duplicates possible

“ Tuples in a bag can have different field lengths and field types

“ Map

$$\left(\text{'alice'}, \left\{ \begin{array}{l} (\text{'lakers'}, 1) \\ (\text{'iPod'}, 2) \end{array} \right\}, [\text{'age'} \rightarrow 20] \right)$$

Data Model

- “ Atom
- “ Tuple
- “ Bag
- “ Map - collection of key-value pairs
 - “ Key is an atom; value can be any type

$$\left(\text{'alice'}, \left\{ \begin{array}{l} (\text{'lakers'}, 1) \\ (\text{'iPod'}, 2) \end{array} \right\}, [\text{'age'} \rightarrow 20] \right)$$

Data Model

“ Control over dataflow

Ex 1 (less efficient)

```
spam_urls = FILTER urls BY isSpam(url);  
culprit_urls = FILTER spam_urls BY pagerank > 0.8;
```

Ex 2 (most efficient)

```
highpgr_urls = FILTER urls BY pagerank > 0.8;  
spam_urls = FILTER highpgr_urls BY isSpam(url);
```

“ Fully nested

- “ More natural for procedural programmers (target user) than normalization

- “ Data is often stored on disk in a nested fashion

- “ Facilitates ease of writing user-defined functions

“ No schema required

Data Model

“ User-Defined Functions (UDFs)

- “ Ex: `spam_urls = FILTER urls BY isSpam(url);`
- “ Can be used in many Pig Latin statements
- “ Useful for custom processing tasks
- “ Can use non-atomic values for input and output
- “ Currently must be written in Java

Speaking Pig Latin

“ LOAD

- “ Input is assumed to be a bag (sequence of tuples)
- “ Can specify a deserializer with “USING”
- “ Can provide a schema with “AS”

```
newBag = LOAD 'filename'  
        <USING functionName() >  
        <AS (fieldName1, fieldName2,...) >;
```

```
Queries = LOAD 'query_log.txt'  
          USING myLoad()  
          AS (userID, queryString, timeStamp)
```


Speaking Pig Latin

“ FOREACH

“ Apply some processing to each tuple in a bag

“ Each field can be:

“ A fieldname of the bag

“ A constant

“ A simple expression (ie: $f1 + f2$)

“ A predefined function (ie: SUM, AVG, COUNT, FLATTEN)

“ A UDF (ie: sumTaxes(gst, pst))

newBag =

FOREACH *bagName*

GENERATE *field1*, *field2*, ...;

Speaking Pig Latin

“ FILTER

“ Select a subset of the tuples in a bag

```
newBag = FILTER bagName  
          BY expression;
```

“ Expression uses simple comparison operators (==, !=, <, >, ...) and Logical connectors (AND, NOT, OR)

```
some_apples =  
  FILTER apples BY colour != 'red';
```

“ Can use UDFs

```
some_apples =  
  FILTER apples BY NOT isRed(colour);
```

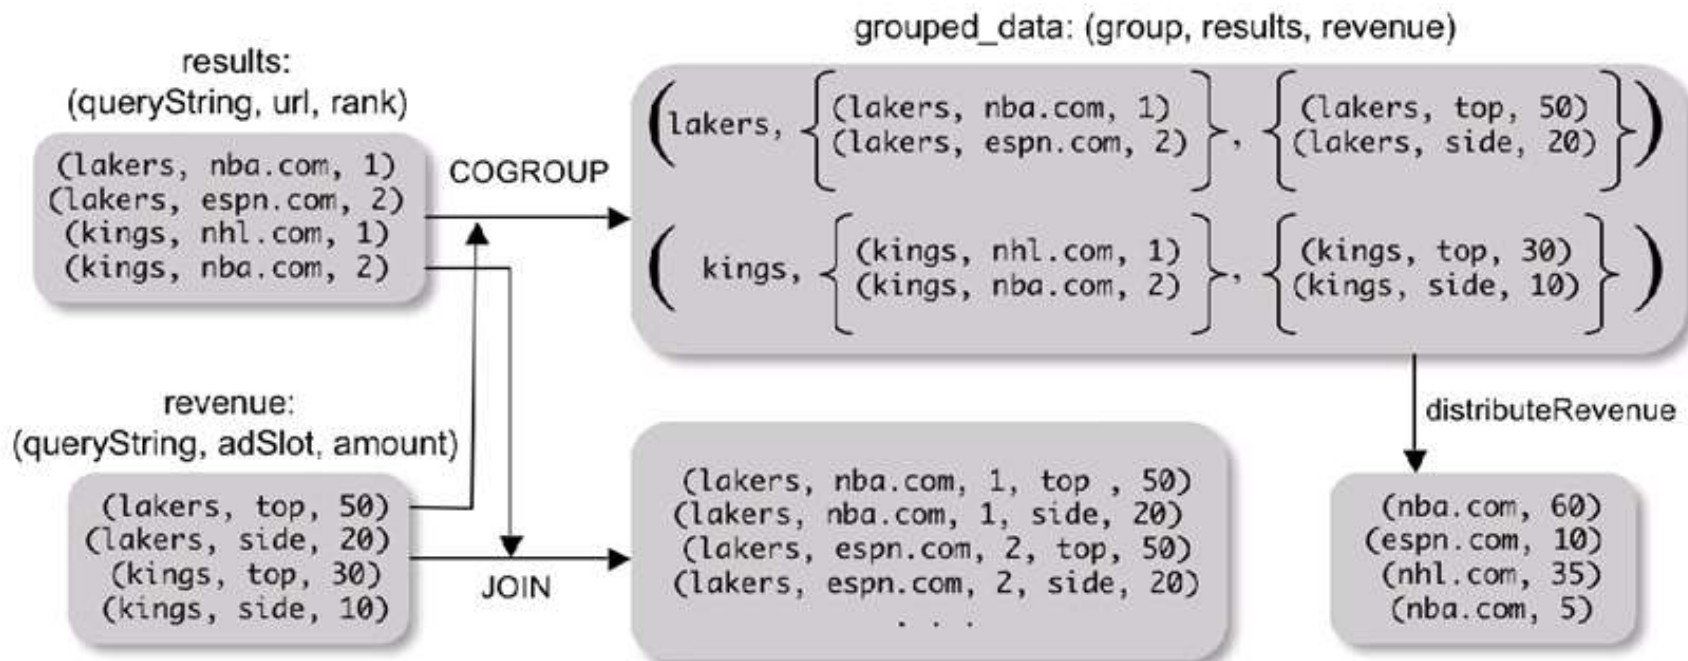
Speaking Pig Latin

“ COGROUP

“ Group two datasets together by a common attribute

“ Groups data into nested bags

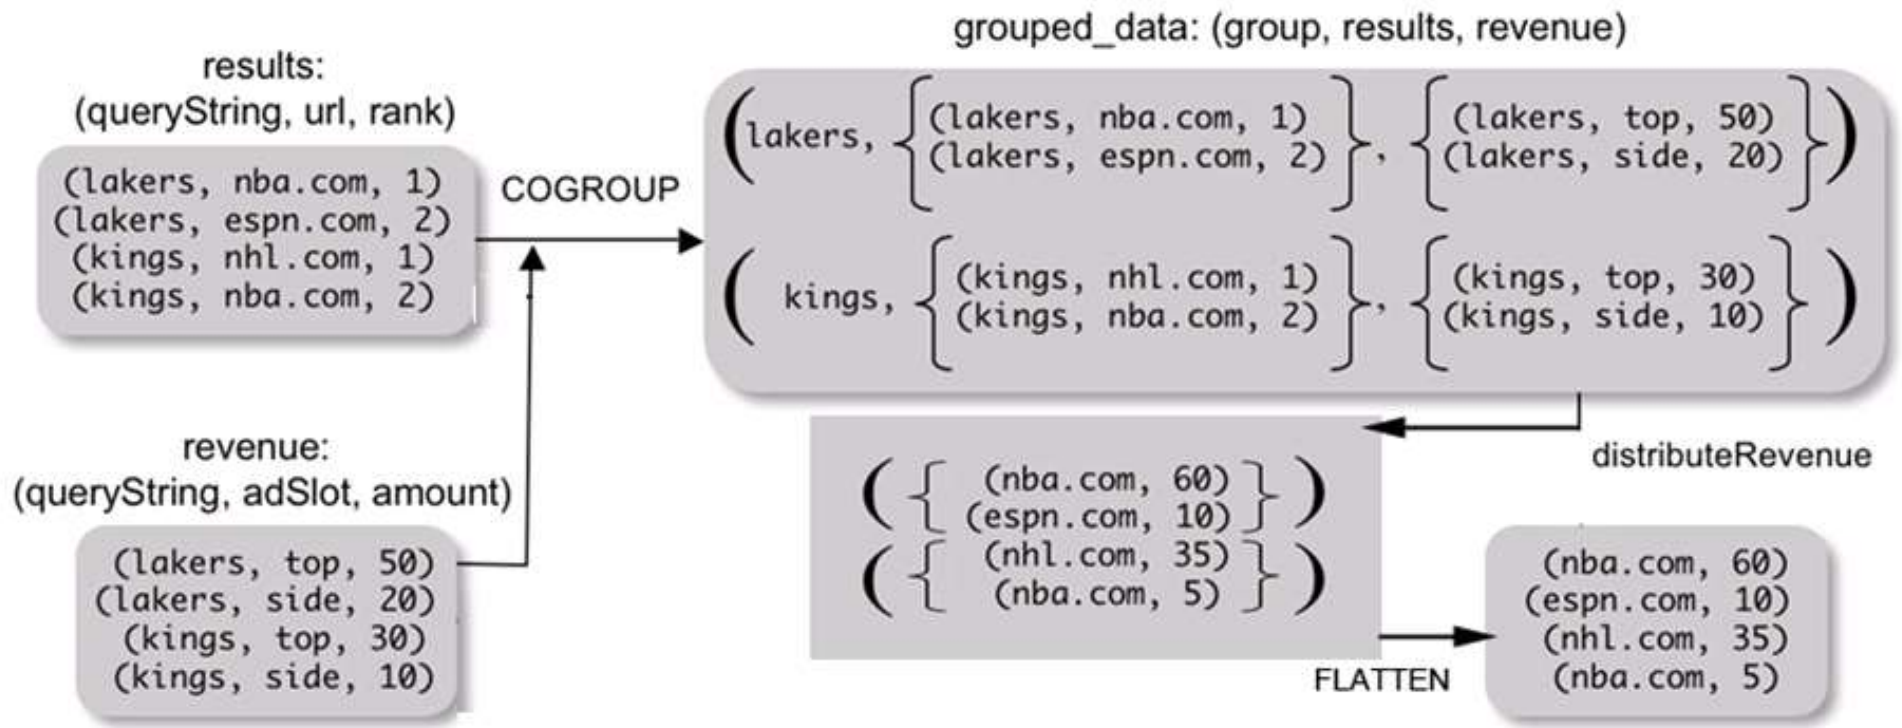
```
grouped_data = COGROUP results BY queryString,  
                           revenue BY queryString;
```



Speaking Pig Latin

“Why COGROUP and not JOIN?”

```
url_revenues =  
    FOREACH grouped_data GENERATE  
    FLATTEN(distributeRev(results, revenue));
```



Speaking Pig Latin

“Why COGROUP and not JOIN?”

- “ May want to process nested bags of tuples before taking the cross product.
- “ Keeps to the goal of a single high-level data transformation per pig-latin statement.
- “ However, JOIN keyword is still available:

```
JOIN results BY queryString,  
     revenue BY queryString;
```



Equivalent

```
temp = COGROUP results BY queryString,  
         revenue BY queryString;  
join_result = FOREACH temp GENERATE  
                FLATTEN(results), FLATTEN(revenue);
```

Speaking Pig Latin

“ STORE (& DUMP)

“ Output data to a file (or screen)

```
STORE bagName INTO 'filename'  
    <USING deserializer ()>;
```

“ Other Commands (incomplete)

“ UNION - return the union of two or more bags

“ CROSS - take the cross product of two or more bags

“ ORDER - order tuples by a specified field(s)

“ DISTINCT - eliminate duplicate tuples in a bag

“ LIMIT - Limit results to a subset

Compilation

- “ Pig system does two tasks:
 - “ Builds a Logical Plan from a Pig Latin script
 - “ Supports execution platform independence
 - “ No processing of data performed at this stage
 - “ Compiles the Logical Plan to a Physical Plan and Executes
 - “ Convert the Logical Plan into a series of Map-Reduce statements to be executed (in this case) by Hadoop Map-Reduce

Compilation

“ Building a Logical Plan

- “ Verify input files and bags referred to are valid
- “ Create a logical plan for each bag(variable) defined

Compilation

“ Building a Logical Plan Example

```
A = LOAD 'user.dat' AS (name, age, city);  
B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```




Load(user.dat)

Compilation

“ Building a Logical Plan Example


```
A = LOAD 'user.dat' AS (name, age, city);  
B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```



Compilation

“ Building a Logical Plan Example


```
A = LOAD 'user.dat' AS (name, age, city);  
B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```



Compilation

“ Building a Logical Plan Example


```
A = LOAD 'user.dat' AS (name, age, city);  
B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```



Compilation

“ Building a Logical Plan Example

```
A = LOAD 'user.dat' AS (name, age, city);  
-> B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```




Compilation

“ Building a Physical Plan

```
A = LOAD 'user.dat' AS (name, age, city);  
B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```


Only happens when output is
specified by STORE or DUMP



Compilation

“ Building a Physical Plan

“ Step 1: Create a map-reduce job for each COGROUP



Compilation

“ Building a Physical Plan


“ Step 1: Create a map-reduce job for each COGROUP

“ Step 2: Push other commands into the map and reduce functions where possible

“ May be the case certain commands require their own map-reduce job (ie: ORDER needs separate map-reduce jobs)

Map

Reduce



Compilation

- “ Efficiency in Execution

- “ Parallelism

- “ Loading data - Files are loaded from HDFS

- “ Statements are compiled into map-reduce jobs

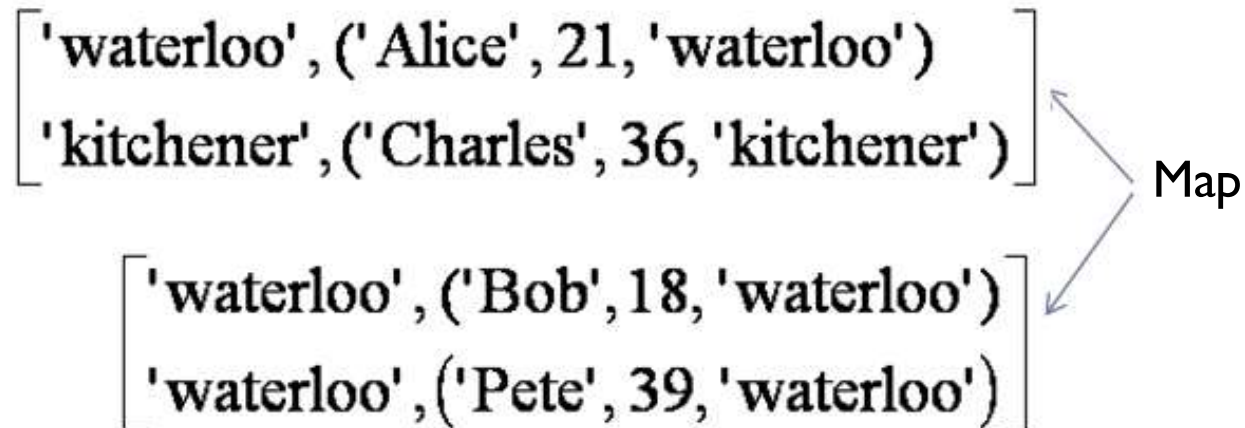
Compilation

“ Efficiency with Nested Bags

- “ In many cases, the nested bags created in each tuple of a COGROUP statement never need to physically materialize
- “ Generally perform aggregation after a COGROUP and the statements for said aggregation are pushed into the reduce function
- “ Applies to algebraic functions (ie: COUNT, MAX, MIN, SUM, AVG)

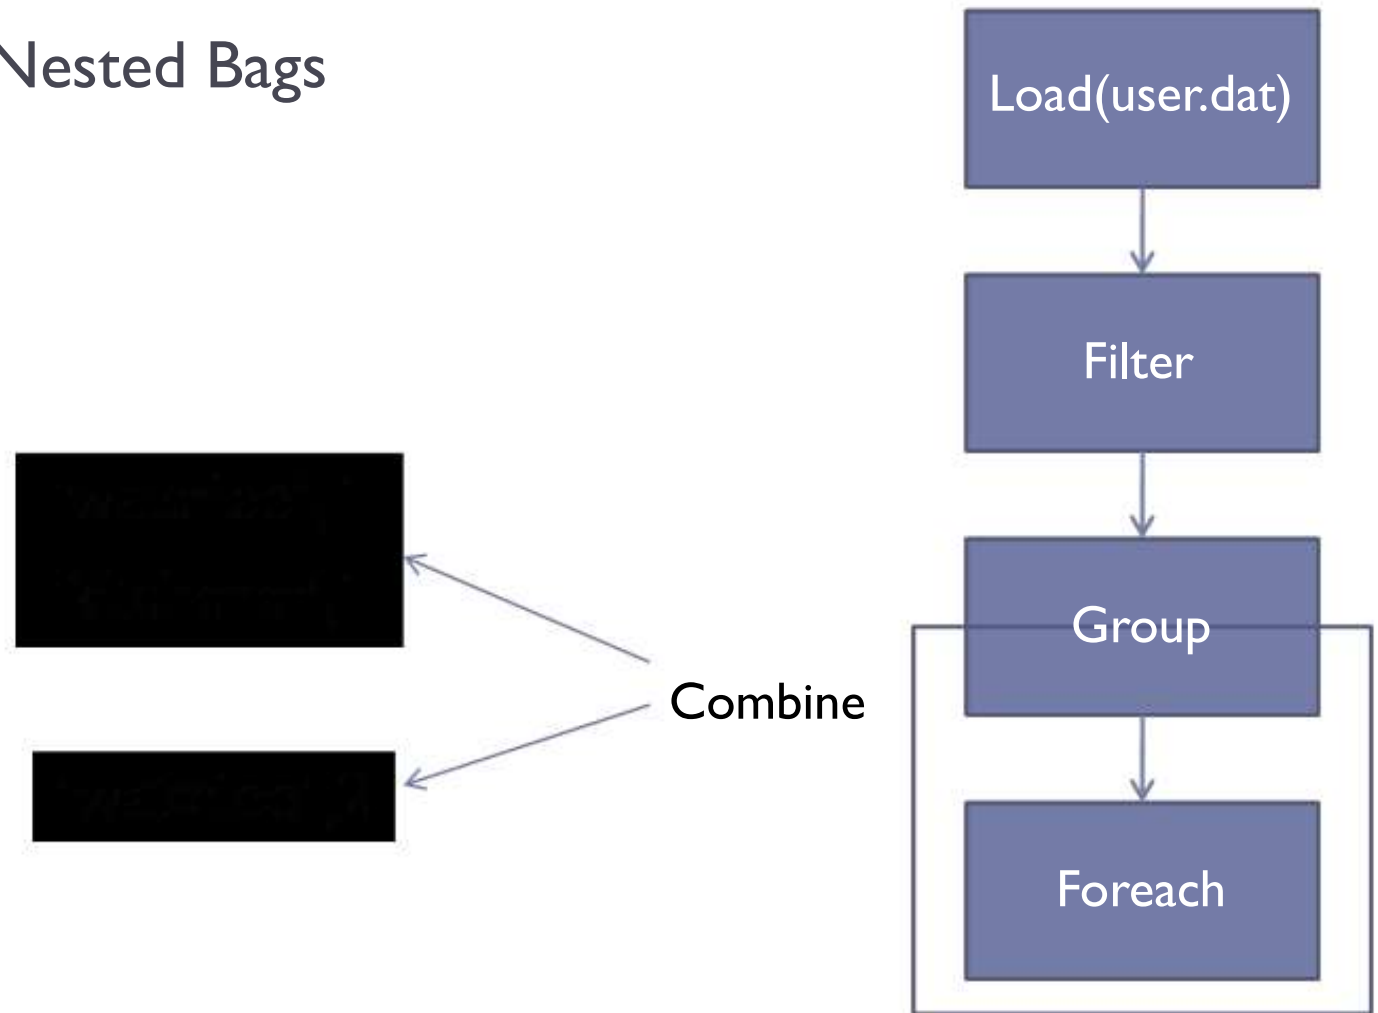
Compilation

“ Efficiency with Nested Bags



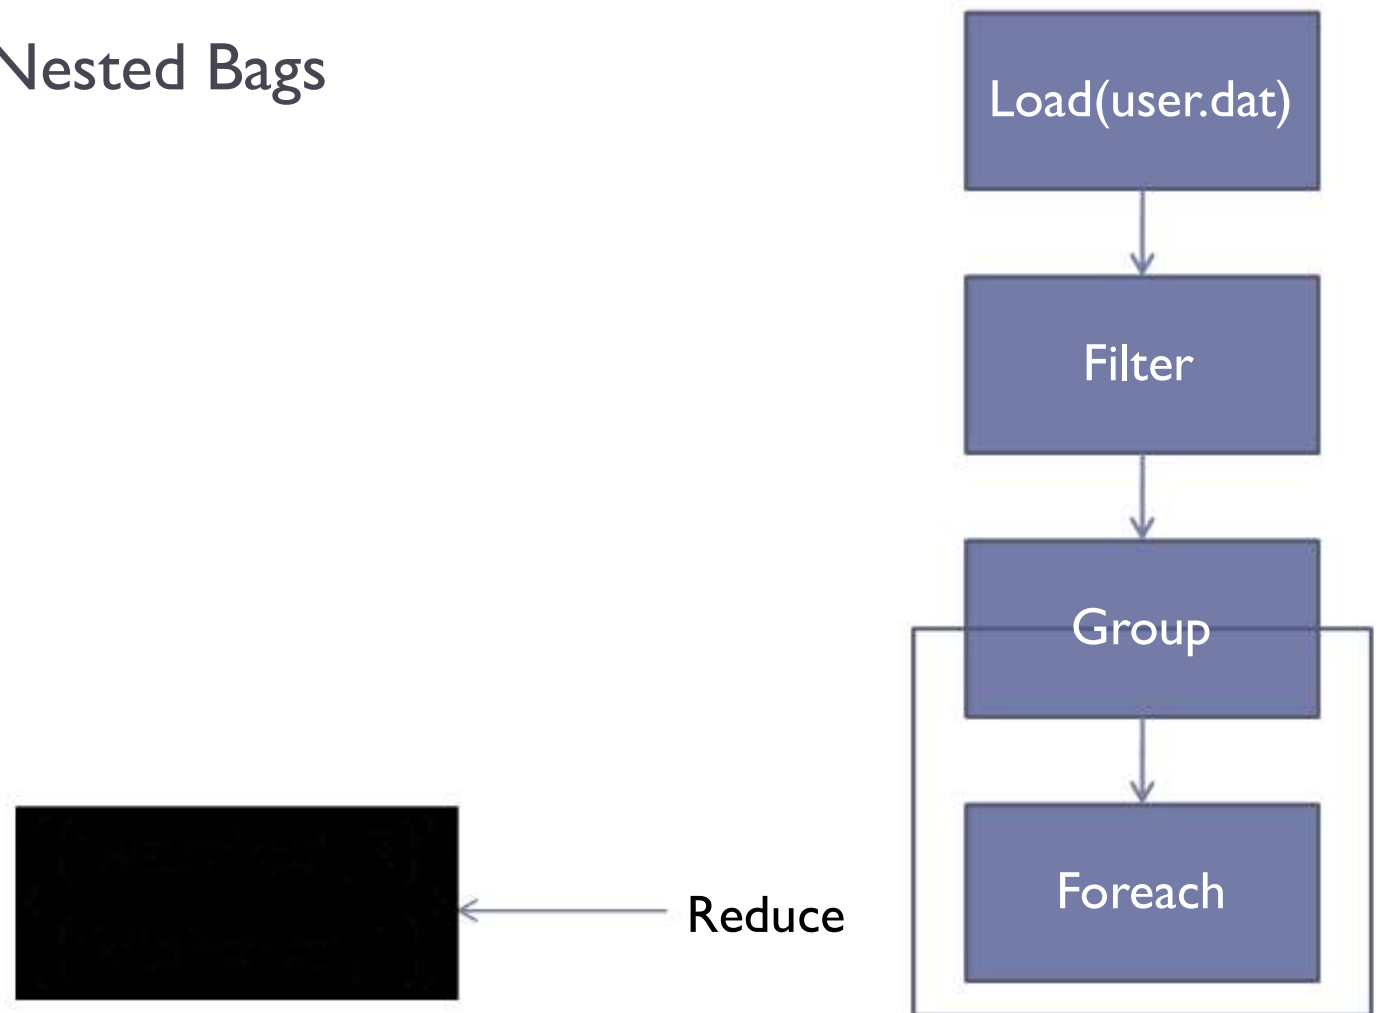
Compilation

“ Efficiency with Nested Bags



Compilation

“ Efficiency with Nested Bags




Compilation

“ Efficiency with Nested Bags

“ Why this works:

- COUNT is an algebraic function; it can be structured as a tree of sub-functions with each leaf working on a subset of the data



Compilation

“ Efficiency with Nested Bags

“ Pig provides an interface for writing algebraic UDFs so they can take advantage of this optimization as well.

“ Inefficiencies

“ Non-algebraic aggregate functions (ie: MEDIAN) need entire bag to materialize; may cause a very large bag to spill to disk if it doesn't fit in memory

“ Every map-reduce job requires data be written and replicated to the HDFS (although this is offset by parallelism achieved)

Debugging

- “ How to verify the semantics of an analysis program
 - “ Run the program against whole data set. Might take hours!
 - “ Generate sample dataset
 - “ Empty result set may occur on few operations like join, filter
 - “ Generally, testing with sample dataset is difficult
- “ Pig-Pen
 - “ Samples data from large dataset for Pig statements
 - “ Apply individual Pig-Latin commands against the dataset
 - “ In case of empty result, pig system resamples
 - “ Remove redundant samples

Debugging

“ Pig-Pen

Operators

LOADGROUPCOGROUPFILTERFOREACHORDER

= LOAD

USING

Default

AS (

)

[Generate Query](#)

<pre>visits = LOAD 'visits.txt' AS (user, url, time); pages = LOAD 'pages.txt' AS (url, pagerank); v_p = JOIN visits BY url, pages BY url; users = GROUP v_p BY user; useravg = FOREACH users GENERATE group, AVG(v_p.pagerank) AS avgpr; answer = FILTER useravg BY avgpr > '0.5';</pre>	<pre>visits: (Amy, cnn.com, 8am) (Amy, frogs.com, 9am) (Fred, snails.com, 11am) pages: (cnn.com, 0.8) (frogs.com, 0.8) (snails.com, 0.3) v_p: (Amy, cnn.com, 8am, cnn.com, 0.8) (Amy, frogs.com, 9am, frogs.com, 0.8) (Fred, snails.com, 11am, snails.com, 0.3) users: (Amy, { (Amy, cnn.com, 8am, cnn.com, 0.8), (Amy, frogs.com, 9am, frogs.com, 0.8) }) (Fred, { (Fred, snails.com, 11am, snails.com, 0.3) }) useravg: (Amy, 0.8) (Fred, 0.3) answer: (Amy, 0.8)</pre>
---	---

Debugging

“ Pig-Latin command window and command generator

The screenshot shows a web-based interface for Pig-Latin. At the top, there's a section titled "Operators" with buttons for LOAD, GROUP, COGROUP, FILTER, FOREACH, and ORDER. Below this is a form to generate a Pig Latin command: `[] = LOAD [] USING [Default] AS ([])`. A "Generate Query" link is also present. The main area is divided into two panes. The left pane contains a Pig Latin script, and the right pane shows the output of that script.

```
visits = LOAD 'visits.txt' AS (user, url, time);

pages = LOAD 'pages.txt' AS (url, pagerank);

v_p = JOIN visits BY url, pages BY url;

users = GROUP v_p BY user;

useravg = FOREACH users GENERATE group, AVG(v_p.pagerank) AS avgpr;

answer = FILTER useravg BY avgpr > '0.5';
```

The output on the right shows the results of each step:

```
visits: (Amy, cnn.com, 8am)
        (Amy, frogs.com, 9am)
        (Fred, snails.com, 11am)

pages: (cnn.com, 0.8)
        (frogs.com, 0.8)
        (snails.com, 0.3)

v_p: (Amy, cnn.com, 8am, cnn.com, 0.8)
      (Amy, frogs.com, 9am, frogs.com, 0.8)
      (Fred, snails.com, 11am, snails.com, 0.3)

users: (Amy, { (Amy, cnn.com, 8am, cnn.com, 0.8),
              (Amy, frogs.com, 9am, frogs.com, 0.8) })
        (Fred, { (Fred, snails.com, 11am, snails.com, 0.3) })

useravg: (Amy, 0.8)
          (Fred, 0.3)

answer: (Amy, 0.8)
```

Debugging

“ Sand Box Dataset (generated automatically!)

Operators

LOADGROUPCOGROUPFILTERFOREACHORDER

= LOAD

USING Default

AS ()

[Generate Query](#)

```
visits = LOAD 'visits.txt' AS (user, url, time);

pages = LOAD 'pages.txt' AS (url, pagerank);

v_p = JOIN visits BY url, pages BY url;

users = GROUP v_p BY user;

useravg = FOREACH users GENERATE group, AVG(v_p.pagerank) AS avgpr;

answer = FILTER useravg BY avgpr > '0.5';
```

```
visits: (Amy, cnn.com, 8am)
        (Amy, frogs.com, 9am)
        (Fred, snails.com, 11am)

pages: (cnn.com, 0.8)
        (frogs.com, 0.8)
        (snails.com, 0.3)

v_p: (Amy, cnn.com, 8am, cnn.com, 0.8)
      (Amy, frogs.com, 9am, frogs.com, 0.8)
      (Fred, snails.com, 11am, snails.com, 0.3)

users: (Amy, { (Amy, cnn.com, 8am, cnn.com, 0.8),
               (Amy, frogs.com, 9am, frogs.com, 0.8) })
        (Fred, { (Fred, snails.com, 11am, snails.com, 0.3) })

useravg: (Amy, 0.8)
          (Fred, 0.3)

answer: (Amy, 0.8)
```

Debugging

“ Pig-Pen

“ Provides sample data that is:

“ Real - taken from actual data

“ Concise - as small as possible

“ Complete - collectively illustrate the key semantics of each command

“ Helps with schema definition

“ Facilitates incremental program writing

Conclusion

“ Pig is a data processing environment in Hadoop that is specifically targeted towards procedural programmers who perform large-scale data analysis.

“ Pig-Latin offers high-level data manipulation in a procedural style.

“ Pig-Pen is a debugging environment for Pig-Latin commands that generates samples from real data.

More Info

“ Pig, <http://hadoop.apache.org/pig/>

“ Hadoop, <http://hadoop.apache.org>

